

▼ Homework 10 - CIFAR10 Image Classification with PyTorch

▼ About

The goal of the homework is to train a convolutional neural network on the standard CIFAR10 image classification dataset.

When solving machine learning tasks using neural networks, one typically starts with a simple network architecture and then improves the network by adding new layers, retraining, adjusting parameters, retraining, etc. We attempt to illustrate this process below with several architecture improvements.

▼ Dev Environment

Working on Google Colab

You may choose to work locally or on Google Colaboratory. You have access to free compute through this service. Colab is recommended since it will be setup correctly and will have access to GPU resources.

1. Visit <https://colab.research.google.com/drive>
2. Navigate to the **Upload** tab, and upload your HW10.ipynb
3. Now on the top right corner, under the Comment and Share options, you should see a Connect option. Once you are connected, you will have access to a VM with 12GB RAM, 50 GB disk space and a single GPU. The dropdown menu will allow you to connect to a local runtime as well.

Notes:

- **If you do not have a working setup for Python 3, this is your best bet. It will also save you from heavy installations like tensorflow if you don't want to deal with those.**
- ***There is a downside.* You can only use this instance for a single 12-hour stretch, after which your data will be deleted, and you would have to redownload all your datasets, any libraries not already on the VM, and regenerate your logs.**

Installing PyTorch and Dependencies

The instructions for installing and setting up PyTorch can be found at <https://pytorch.org/get-started/locally/>. Make sure you follow the instructions for your machine. For any of the remaining libraries used in this assignment:

- We have provided a `hw8_requirements.txt` file on the homework web page.
- Download this file, and in the same directory you can run `pip3 install -r hw8_requirements.txt` Check that PyTorch installed correctly by running the following:

```
1 import torch  
2 torch.rand(5, 3)  
  
↳ tensor([[0.1072, 0.1348, 0.5591],  
         [0.0972, 0.1583, 0.9011],  
         [0.3416, 0.6958, 0.7808],  
         [0.5589, 0.8751, 0.1891],  
         [0.6782, 0.3040, 0.4816]])
```

▼ Part 0 Imports and Basic Setup (5 Points)

First, import the required libraries as follows. The libraries we will use will be the same as those in HW8.

```
1 import numpy as np  
2 import torch  
3 from torch import nn  
4 from torch import optim  
5 from torchsummary import summary  
6  
7 import matplotlib.pyplot as plt
```

GPU Support

Training of large network can take a long time. PyTorch supports GPU with just a small amount of effort.

When creating our networks, we will call `net.to(device)` to tell the network to train on the GPU, if one is available. Note, if the network utilizes the GPU, it is important that any tensors we use with it (such as the data) also reside on the CPU. Thus, a call like `images = images.to(device)` is necessary with any data we want to use with the GPU.

Note: If you can't get access to a GPU, don't worry too much. Since we use very small networks, the difference between CPU and GPU isn't large and in some cases GPU will actually be slower.

```
1 import torch.cuda as cuda
```

```
2  
3 # Use a GPU, i.e. cuda:0 device if it available.  
4 device = torch.device("cuda:0" if cuda.is_available() else "cpu")  
5 print(device)
```

```
↪ cuda:0
```

▼ Training Code

```
1 import time  
2  
3 class Flatten(nn.Module):  
4     """NN Module that flattens the incoming tensor."""  
5     def forward(self, input):  
6         return input.view(input.size(0), -1)  
7  
8 def train(model, train_loader, test_loader, loss_func, opt, num_epochs=10):  
9     all_training_loss = np.zeros((0,2))  
10    all_training_acc = np.zeros((0,2))  
11    all_test_loss = np.zeros((0,2))  
12    all_test_acc = np.zeros((0,2))  
13  
14    training_step = 0  
15    training_loss, training_acc = 2.0, 0.0  
16    print_every = 1000  
17  
18    start = time.clock()  
19  
20    for i in range(num_epochs):  
21        epoch_start = time.clock()  
22  
23        model.train()  
24        for images, labels in train_loader:  
25            images, labels = images.to(device), labels.to(device)  
26            opt.zero_grad()  
27  
28            preds = model(images)  
29            loss = loss_func(preds, labels)  
30            loss.backward()  
31            opt.step()  
32  
33            training_loss += loss.item()  
34            training_acc += (torch.argmax(preds, dim=1)==labels).float().mean()  
35  
36        if training_step % print_every == 0:  
37            training_loss /= print_every
```

```
38     training_acc /= print_every
39
40     all_training_loss = np.concatenate((all_training_loss, [[training_step, training_loss]]))
41     all_training_acc = np.concatenate((all_training_acc, [[training_step, training_acc]]))
42
43     print(' Epoch %d @ step %d: Train Loss: %3f, Train Accuracy: %3f' % (
44         i, training_step, training_loss, training_acc))
45     training_loss, training_acc = 0.0, 0.0
46
47     training_step+=1
48
49 model.eval()
50 with torch.no_grad():
51     validation_loss, validation_acc = 0.0, 0.0
52     count = 0
53     for images, labels in test_loader:
54         images, labels = images.to(device), labels.to(device)
55         output = model(images)
56         validation_loss+=loss_func(output,labels)
57         validation_acc+=(torch.argmax(output, dim=1) == labels).float().mean()
58         count += 1
59     validation_loss/=count
60     validation_acc/=count
61
62     all_test_loss = np.concatenate((all_test_loss, [[training_step, validation_loss]]))
63     all_test_acc = np.concatenate((all_test_acc, [[training_step, validation_acc]]))
64
65     epoch_time = time.clock() - epoch_start
66
67     print('Epoch %d Test Loss: %3f, Test Accuracy: %3f, time: %.1fs' % (
68         i, validation_loss, validation_acc, epoch_time))
69
70     total_time = time.clock() - start
71     print('Final Test Loss: %3f, Test Accuracy: %3f, Total time: %.1fs' % (
72         validation_loss, validation_acc, total_time))
73
74     return {'loss': { 'train': all_training_loss, 'test': all_test_loss },
75             'accuracy': { 'train': all_training_acc, 'test': all_test_acc }}
76
77 def plot_graphs(model_name, metrics):
78     for metric, values in metrics.items():
79         for name, v in values.items():
80             plt.plot(v[:,0], v[:,1], label=name)
81     plt.title(f'{metric} for {model_name}')
82     plt.legend()
83     plt.xlabel("Training Steps")
84     plt.ylabel(metric)
85     plt.show()
86
```

Load the **CIFAR-10** dataset and define the transformations. You may also want to print its structure, size, as well as sample a few images to get a sense of how to design the network.

```
1 !mkdir hw10_data

1 # Download the data.
2 from torchvision import datasets, transforms
3
4 transformations = transforms.Compose(
5     [transforms.ToTensor(),
6      transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
7 train_set = datasets.CIFAR10(root='hw10_data/', download=True, transform=transformations)
8 test_set = datasets.CIFAR10(root='hw10_data', download=True, train=False, transform=transformations)
```

```
↳ 0% | 0/170498071 [00:00<?, ?it/s] Downloading https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz to hw10_dat
170500096it [00:07, 23606943.09it/s]
Files already downloaded and verified
```

Use DataLoader to create a loader for the training set and a loader for the testing set. You can use a batch_size of 8 to start, and change it if you wish.

```
1 from torch.utils.data import DataLoader
2
3 batch_size = 8
4 train_loader = torch.utils.data.DataLoader(train_set, batch_size, shuffle=True, num_workers=2)
5 test_loader = torch.utils.data.DataLoader(test_set, batch_size, shuffle=True, num_workers=2)
6
7 input_shape = np.array(train_set[0][0]).shape
8 input_dim = input_shape[1]*input_shape[2]*input_shape[0]
9

1 training_epochs = 5
```

▼ Part 1 CIFAR10 with Fully Connected Neural Netowrk (25 Points)

As a warm-up, let's begin by training a two-layer fully connected neural network model on **CIFAR-10** dataset. You may go back to check HW8 for some basics.

We will give you this code to use as a baseline to compare against your CNN models.

```
1 class TwoLayerModel(nn.Module):
2     def __init__(self):
3         super(TwoLayerModel, self).__init__()
4         self.net = nn.Sequential(
5             Flatten(),
6             nn.Linear(input_dim, 64),
7             nn.ReLU(),
8             nn.Linear(64, 10))
9
10    def forward(self, x):
11        return self.net(x)
12
13 model = TwoLayerModel().to(device)
14
15 loss = nn.CrossEntropyLoss()
16 optimizer = optim.RMSprop(model.parameters(), lr=0.001, weight_decay=0.01)
17
18 # Training epoch should be about 15-20 sec each on GPU.
19 metrics = train(model, train_loader, test_loader, loss, optimizer, training_epochs)
```



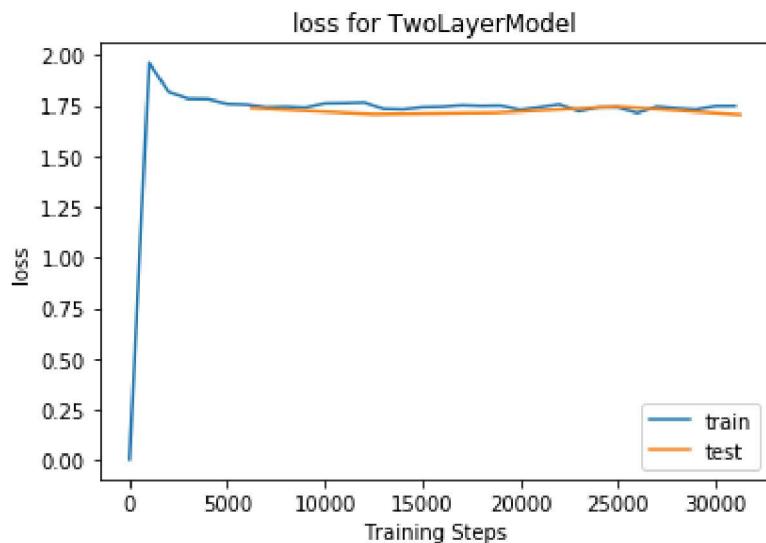
```
Epoch 0 @ step 0: Train Loss: 0.004301, Train Accuracy: 0.000125
Epoch 0 @ step 1000: Train Loss: 1.961467, Train Accuracy: 0.320250
Epoch 0 @ step 2000: Train Loss: 1.817933, Train Accuracy: 0.352000
Epoch 0 @ step 3000: Train Loss: 1.783904, Train Accuracy: 0.357500
Epoch 0 @ step 4000: Train Loss: 1.782436, Train Accuracy: 0.366000
Epoch 0 @ step 5000: Train Loss: 1.758631, Train Accuracy: 0.374250
Epoch 0 @ step 6000: Train Loss: 1.755314, Train Accuracy: 0.368750
Epoch 0 Test Loss: 1.738084, Test Accuracy: 0.378600, time: 16.6s
Epoch 1 @ step 7000: Train Loss: 1.742370, Train Accuracy: 0.383625
Epoch 1 @ step 8000: Train Loss: 1.744828, Train Accuracy: 0.374375
Epoch 1 @ step 9000: Train Loss: 1.740492, Train Accuracy: 0.374625
Epoch 1 @ step 10000: Train Loss: 1.761740, Train Accuracy: 0.368875
Epoch 1 @ step 11000: Train Loss: 1.763021, Train Accuracy: 0.372375
Epoch 1 @ step 12000: Train Loss: 1.766127, Train Accuracy: 0.367250
Epoch 1 Test Loss: 1.707702, Test Accuracy: 0.386500, time: 17.7s
Epoch 2 @ step 13000: Train Loss: 1.734238, Train Accuracy: 0.381125
```

Plot the model results

Normally we would want to use Tensorboard for looking at metrics. However, if colab reset while we are working, we might lose our logs and therefore our metrics. Let's just plot some graphs that will survive across colab instances.

```
Epoch 2 Test Loss: 1.715050  Test Accuracy: 0.387000  time: 16.8s
1| plot_graphs("TwoLayerModel", metrics)
```





▼ Part 2 Convolutional Neural Network (CNN) (35 Points)

Now, let's design a convolution neural netwrok!

Build a simple CNN model, inserting 2 CNN layers in front of our 2 layer fully connect model from above:

1. A convolution with 3x3 filter, 16 output channels, stride = 1, padding=1
2. A ReLU activation
3. A Max-Pooling layer with 2x2 window
4. A convolution, 3x3 filter, 16 output channels, stride = 1, padding=1
5. A ReLU activation
6. Flatten layer
7. Fully connected linear layer with output size 64
8. ReLU
9. Fully connected linear layer, with output size 10

You will have to figure out the input sizes of the first fully connected layer based on the previous layer sizes. Note that you also need to fill those in the report section (see report section in the notebook for details)

```
1 | layer1_in_channels = 3
2 | layer1_output_channels = 16
3 |
```

```
4 layer2_in_channels = layer1_output_channels
5 layer2_output_channels = 16
6
7 layer3_out = 64
8 kernel_size = 3
9 nun_labels = 10
10
11 padding_size=1
12 stride_size=1
13
14 class ConvModel(nn.Module):
15     # Your Code Here
16     def __init__(self):
17         super().__init__()
18         self.layer1 = nn.Sequential(
19             nn.Conv2d(layer1_in_channels, layer1_output_channels, kernel_size, stride=stride_size, padding=padding_size),
20             nn.ReLU(),
21             nn.MaxPool2d(2))
22
23         self.layer2 = nn.Sequential(
24             nn.Conv2d(layer2_in_channels, layer2_output_channels, kernel_size, stride=stride_size, padding=padding_size),
25             nn.ReLU(),
26             Flatten())
27
28
29         self.layer3 = nn.Sequential(
30             nn.Linear(4096, layer3_out),
31             nn.ReLU())
32
33         self.fc = nn.Linear(layer3_out, nun_labels)
34
35
36     def forward(self, x):
37         out = self.layer1(x)
38         out = self.layer2(out)
39         out = self.layer3(out)
40         out = out.reshape(out.size(0), -1)
41         return self.fc(out)
42
43
44
45 print(ConvModel())
46
47 model = ConvModel().to(device)
48
49 summary(model, (3,32,32))
50
51 loss = nn.CrossEntropyLoss()
52 optimizer = optim.RMSprop(model.parameters(), lr=0.001, weight_decay=0.01)
53
54 metrics = train(model, train_loader, test_loader, loss, optimizer, training_epochs)
```

55



```
ConvModel(  
    (layer1): Sequential(  
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (1): ReLU()  
        (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    )  
    (layer2): Sequential(  
        (0): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (1): ReLU()  
        (2): Flatten()  
    )  
    (layer3): Sequential(  
        (0): Linear(in_features=4096, out_features=64, bias=True)  
        (1): ReLU()  
    )  
    (fc): Linear(in_features=64, out_features=10, bias=True)  
)
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 16, 32, 32]	448
ReLU-2	[-1, 16, 32, 32]	0
MaxPool2d-3	[-1, 16, 16, 16]	0
Conv2d-4	[-1, 16, 16, 16]	2,320
ReLU-5	[-1, 16, 16, 16]	0
Flatten-6	[-1, 4096]	0
Linear-7	[-1, 64]	262,208
ReLU-8	[-1, 64]	0
Linear-9	[-1, 10]	650

Total params: 265,626

Trainable params: 265,626

Non-trainable params: 0

Input size (MB): 0.01

Forward/backward pass size (MB): 0.38

Params size (MB): 1.01

Estimated Total Size (MB): 1.40

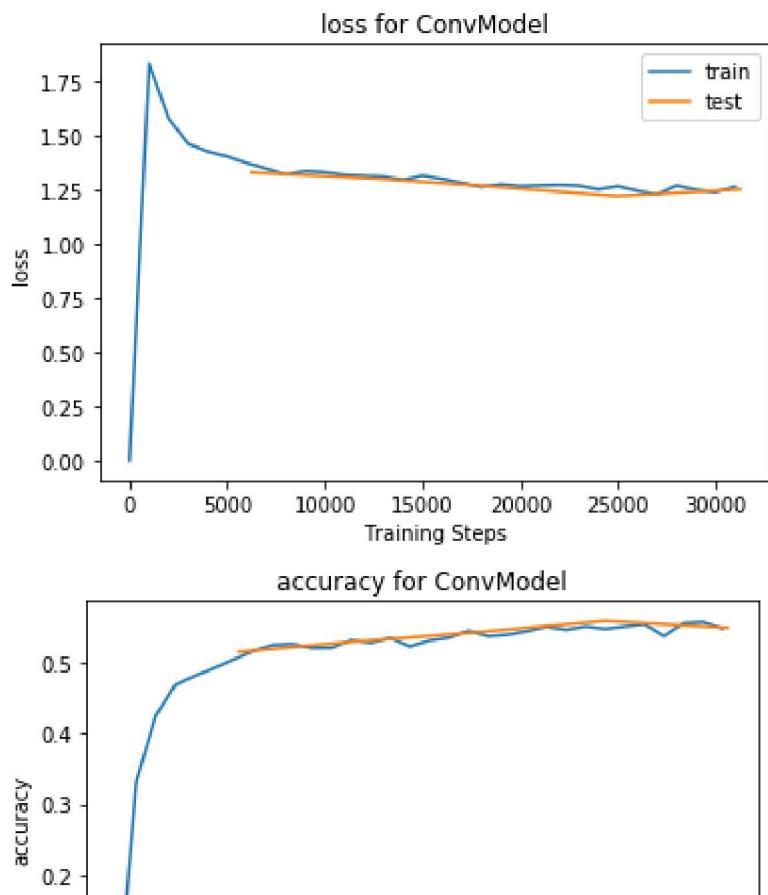
Exception ignored in: <bound method _DataLoaderIter.__del__ of <torch.utils.data.dataloader._DataLoaderIter object at 0x7fe6
Traceback (most recent call last):

File "/usr/local/lib/python3.6/dist-packages/torch/utils/data/dataloader.py". line 717. in del

```
Traceback (most recent call last):
  self._shutdown_workers()
Exception ignored in: <bound method _DataLoaderIter.__del__ of <torch.utils.data.dataloader._DataLoaderIter object at 0x7fe6
Traceback (most recent call last):
  File "/usr/local/lib/python3.6/dist-packages/torch/utils/data/dataloader.py", line 713, in _shutdown_workers
  File "/usr/lib/python3.6/multiprocessing/queues.py", line 240, in _feed
    send_bytes(obj)
    w.join()
  File "/usr/lib/python3.6/multiprocessing/process.py", line 122, in join
  File "/usr/lib/python3.6/multiprocessing/connection.py", line 200, in send_bytes
    self._send_bytes(m[offset:offset + size])
    assert self._parent_pid == os.getpid(), 'can only join a child process'
  File "/usr/lib/python3.6/multiprocessing/connection.py", line 404, in _send_bytes
    self._send(header + buf)
AssertionError: can only join a child process
  File "/usr/lib/python3.6/multiprocessing/queues.py", line 240, in _feed
    send_bytes(obj)
Traceback (most recent call last):
  File "/usr/lib/python3.6/multiprocessing/connection.py", line 368, in _send
    n = write(self._handle, buf)
BrokenPipeError: [Errno 32] Broken pipe
  File "/usr/local/lib/python3.6/dist-packages/torch/utils/data/dataloader.py", line 717, in __del__
  File "/usr/lib/python3.6/multiprocessing/connection.py", line 200, in send_bytes
    self._send_bytes(m[offset:offset + size])
  File "/usr/lib/python3.6/multiprocessing/connection.py", line 404, in _send_bytes
    self._send(header + buf)
    self._shutdown_workers()
  File "/usr/local/lib/python3.6/dist-packages/torch/utils/data/dataloader.py", line 713, in _shutdown_workers
    w.join()
  File "/usr/lib/python3.6/multiprocessing/connection.py", line 368, in _send
    n = write(self._handle, buf)
  File "/usr/lib/python3.6/multiprocessing/process.py", line 122, in join
```

```
1| plot_graphs("ConvModel", metrics)
```





Do you notice the improvement over the accuracy compared to that in Part 1?



▼ Part 3 Open Design Competition (35 Points + 10 bonus points)

Try to beat the previous models by adding additional layers, changing parameters, etc. You should add at least one layer.

Possible changes include:

- Dropout
- Batch Normalization
- More layers
- Residual Connections (harder)

- Change layer size
- Pooling layers, stride
- Different optimizer
- Train for longer

Once you have a model you think is great, evaluate it against our hidden test data (see hidden_loader above) and upload the results to the leader board on gradescope. **The top 3 scorers will get a bonus 10 points.**

You can steal model structures found on the internet if you want. The only constraint is that **you must train the model from scratch**.

```
1 # You Awesome Super Best model code here
2 layer1_in_channels = 3
3 layer1_output_channels = 128
4
5 layer2_in_channels = layer1_output_channels
6 layer2_output_channels = 128
7
8 layer3_in_channels = layer2_output_channels
9 layer3_output_channels = 128
10
11 layer4_in_channels = layer3_output_channels
12 layer4_output_channels = 128
13
14 kernel_size = 3
15 nun_labels = 10
16
17
18 class AwesoneModel(nn.Module):
19     def __init__(self):
20         super().__init__()
21         self.layer1 = nn.Sequential(
22             nn.Conv2d(layer1_in_channels, layer1_output_channels, kernel_size, padding=1),
23             nn.ReLU(),
24             #nn.Conv2d(layer1_output_channels, layer1_output_channels, kernel_size, padding=1),
25             #nn.ReLU(),
26             nn.MaxPool2d(2))
27
28         self.layer2 = nn.Sequential(
29             nn.Conv2d(layer2_in_channels, layer2_output_channels, kernel_size, padding=1),
30             nn.ReLU(),
31             #nn.Conv2d(layer2_output_channels, layer2_output_channels, kernel_size, padding=1),
32             #nn.ReLU(),
33             #nn.Conv2d(layer2_output_channels, layer2_output_channels, kernel_size, padding=1),
34             #nn.ReLU(),
35             nn.MaxPool2d(2))
36
37         self.layer3 = nn.Sequential(
```

```
38     nn.Conv2d(layer3_in_channels, layer3_output_channels, kernel_size),
39     nn.BatchNorm2d(layer3_output_channels),
40     nn.ReLU(),
41     nn.MaxPool2d(2))
42
43     self.dropout = nn.Dropout2d(p=0.2)
44
45     self.layer4 = nn.Sequential(
46         nn.Conv2d(layer4_in_channels, layer4_output_channels, kernel_size, padding=1),
47         nn.BatchNorm2d(layer4_output_channels),
48         nn.ReLU(),
49         nn.MaxPool2d(2))
50
51     self.fc = nn.Linear(layer4_output_channels, num_labels)
52
53 def forward(self, x):
54     out = self.layer1(x)
55
56     out = self.layer2(out)
57
58     out = self.layer3(out)
59
60     out = self.dropout(out)
61
62     out = self.layer4(out)
63
64     out = out.reshape(out.size(0), -1)
65     return self.fc(out)
66
67
68
69
70
71 model = AwesomeModel().to(device)
72
73 loss = nn.CrossEntropyLoss()
74 # optimizer = optim.RMSprop(model.parameters(), lr=0.001, weight_decay=0.01)
75 optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)
76
77 summary(model, (3,32,32))
78 metrics = train(model, train_loader, test_loader, loss, optimizer, 10)
```



Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 128, 32, 32]	3,584
ReLU-2	[-1, 128, 32, 32]	0
MaxPool2d-3	[-1, 128, 16, 16]	0
Conv2d-4	[-1, 128, 16, 16]	147,584
ReLU-5	[-1, 128, 16, 16]	0
MaxPool2d-6	[-1, 128, 8, 8]	0
Conv2d-7	[-1, 128, 6, 6]	147,584
BatchNorm2d-8	[-1, 128, 6, 6]	256
ReLU-9	[-1, 128, 6, 6]	0
MaxPool2d-10	[-1, 128, 3, 3]	0
Dropout2d-11	[-1, 128, 3, 3]	0
Conv2d-12	[-1, 128, 3, 3]	147,584
BatchNorm2d-13	[-1, 128, 3, 3]	256
ReLU-14	[-1, 128, 3, 3]	0
MaxPool2d-15	[-1, 128, 1, 1]	0
Linear-16	[-1, 10]	1,290

Total params: 448,138

Trainable params: 448,138

Non-trainable params: 0

Input size (MB): 0.01

Forward/backward pass size (MB): 2.96

Params size (MB): 1.71

Estimated Total Size (MB): 4.68

Epoch 0 @ step 0: Train Loss: 0.004442, Train Accuracy: 0.000125

Epoch 0 @ step 1000: Train Loss: 1.760698, Train Accuracy: 0.350500

Epoch 0 @ step 2000: Train Loss: 1.466611, Train Accuracy: 0.469375

Epoch 0 @ step 3000: Train Loss: 1.319344, Train Accuracy: 0.527750

Epoch 0 @ step 4000: Train Loss: 1.225576, Train Accuracy: 0.568000

Epoch 0 @ step 5000: Train Loss: 1.152384, Train Accuracy: 0.594125

Epoch 0 @ step 6000: Train Loss: 1.083020, Train Accuracy: 0.613125

Epoch 0 Test Loss: 0.952849, Test Accuracy: 0.659100, time: 31.8s

Epoch 1 @ step 7000: Train Loss: 1.030204, Train Accuracy: 0.640500

Epoch 1 @ step 8000: Train Loss: 0.981465, Train Accuracy: 0.657250

Epoch 1 @ step 9000: Train Loss: 0.962273, Train Accuracy: 0.666125

Epoch 1 @ step 10000: Train Loss: 0.941776, Train Accuracy: 0.672875

Epoch 1 @ step 11000: Train Loss: 0.909894, Train Accuracy: 0.685375

```
--> Epoch 1 @ step 12000: Train Loss: 0.915352, Train Accuracy: 0.677375
Epoch 1 Test Loss: 0.831166, Test Accuracy: 0.716500, time: 33.0s
    Epoch 2 @ step 13000: Train Loss: 0.845838, Train Accuracy: 0.705250
    Epoch 2 @ step 14000: Train Loss: 0.827135, Train Accuracy: 0.712000
    Epoch 2 @ step 15000: Train Loss: 0.807027, Train Accuracy: 0.723375
    Epoch 2 @ step 16000: Train Loss: 0.808400, Train Accuracy: 0.723375
    Epoch 2 @ step 17000: Train Loss: 0.791600, Train Accuracy: 0.723875
    Epoch 2 @ step 18000: Train Loss: 0.798014, Train Accuracy: 0.722125
Epoch 2 Test Loss: 0.754116, Test Accuracy: 0.739400, time: 32.4s
    Epoch 3 @ step 19000: Train Loss: 0.749551, Train Accuracy: 0.741000
    Epoch 3 @ step 20000: Train Loss: 0.698005, Train Accuracy: 0.754625
    Epoch 3 @ step 21000: Train Loss: 0.717129, Train Accuracy: 0.745750
    Epoch 3 @ step 22000: Train Loss: 0.699480, Train Accuracy: 0.757875
    Epoch 3 @ step 23000: Train Loss: 0.703705, Train Accuracy: 0.753625
    Epoch 3 @ step 24000: Train Loss: 0.702818, Train Accuracy: 0.757000
Epoch 3 Test Loss: 0.686003, Test Accuracy: 0.763200, time: 32.4s
    Epoch 4 @ step 25000: Train Loss: 0.697754, Train Accuracy: 0.757625
    Epoch 4 @ step 26000: Train Loss: 0.613823, Train Accuracy: 0.783250
    Epoch 4 @ step 27000: Train Loss: 0.607760, Train Accuracy: 0.787625
    Epoch 4 @ step 28000: Train Loss: 0.629644, Train Accuracy: 0.783375
    Epoch 4 @ step 29000: Train Loss: 0.618953, Train Accuracy: 0.781750
    Epoch 4 @ step 30000: Train Loss: 0.626694, Train Accuracy: 0.780875
    Epoch 4 @ step 31000: Train Loss: 0.635129, Train Accuracy: 0.774750
Epoch 4 Test Loss: 0.673252, Test Accuracy: 0.767200, time: 32.3s
    Epoch 5 @ step 32000: Train Loss: 0.559631, Train Accuracy: 0.806500
    Epoch 5 @ step 33000: Train Loss: 0.550793, Train Accuracy: 0.810750
    Epoch 5 @ step 34000: Train Loss: 0.548881, Train Accuracy: 0.808500
    Epoch 5 @ step 35000: Train Loss: 0.565127, Train Accuracy: 0.801000
    Epoch 5 @ step 36000: Train Loss: 0.557797, Train Accuracy: 0.806125
    Epoch 5 @ step 37000: Train Loss: 0.577192, Train Accuracy: 0.798625
Epoch 5 Test Loss: 0.643921, Test Accuracy: 0.780900, time: 32.1s
    Epoch 6 @ step 38000: Train Loss: 0.496895, Train Accuracy: 0.821250
    Epoch 6 @ step 39000: Train Loss: 0.488642, Train Accuracy: 0.836000
    Epoch 6 @ step 40000: Train Loss: 0.510050, Train Accuracy: 0.820750
    Epoch 6 @ step 41000: Train Loss: 0.512778, Train Accuracy: 0.821250
    Epoch 6 @ step 42000: Train Loss: 0.528616, Train Accuracy: 0.816250
    Epoch 6 @ step 43000: Train Loss: 0.531179, Train Accuracy: 0.812250
Epoch 6 Test Loss: 0.603606, Test Accuracy: 0.796800, time: 32.8s
    Epoch 7 @ step 44000: Train Loss: 0.496106, Train Accuracy: 0.827375
    Epoch 7 @ step 45000: Train Loss: 0.443340, Train Accuracy: 0.844500
    Epoch 7 @ step 46000: Train Loss: 0.431364, Train Accuracy: 0.850875
    Epoch 7 @ step 47000: Train Loss: 0.464086, Train Accuracy: 0.837125
```

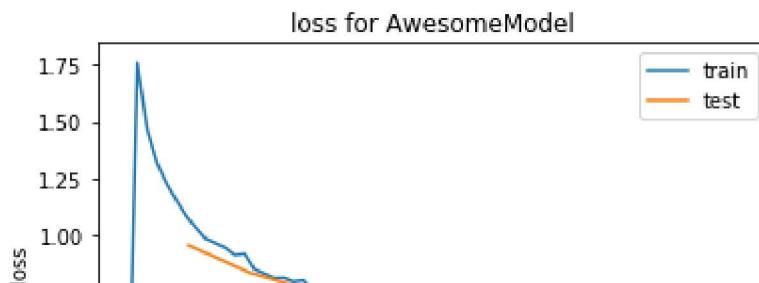
```
Epoch 7 @ step 48000: Train Loss: 0.470960, Train Accuracy: 0.836750
Epoch 7 @ step 49000: Train Loss: 0.476336, Train Accuracy: 0.832500
Epoch 7 Test Loss: 0.616699, Test Accuracy: 0.798900, time: 32.6s
Epoch 8 @ step 50000: Train Loss: 0.453773, Train Accuracy: 0.838500
Epoch 8 @ step 51000: Train Loss: 0.391436, Train Accuracy: 0.860125
Epoch 8 @ step 52000: Train Loss: 0.423008, Train Accuracy: 0.849625
Epoch 8 @ step 53000: Train Loss: 0.406135, Train Accuracy: 0.859375
Epoch 8 @ step 54000: Train Loss: 0.401967, Train Accuracy: 0.858375
Epoch 8 @ step 55000: Train Loss: 0.436225, Train Accuracy: 0.847875
Epoch 8 @ step 56000: Train Loss: 0.450351, Train Accuracy: 0.841125
Epoch 8 Test Loss: 0.615750, Test Accuracy: 0.796400, time: 33.0s
```

What changes did you make to improve your model?

```
Epoch 7 @ step 50000: Train Loss: 0.391436, Train Accuracy: 0.860125
```

```
1| plot_graphs("AwesomeModel", metrics)
```





After you get a nice model, download the test_file.zip and unzip it to get test_file.pt. In colab, you can explore your files from the left side bar. You can also download the files to your machine from there.

```
0.25 1 |  
|  
1 !wget http://courses.engr.illinois.edu/cs498aml/sp2019/homeworks/test_file.zip  
2 !unzip test_file.zip  
  
[  --2019-05-02 04:53:55-- http://courses.engr.illinois.edu/cs498aml/sp2019/homeworks/test_file.zip  
Resolving courses.engr.illinois.edu (courses.engr.illinois.edu)... 130.126.151.9  
Connecting to courses.engr.illinois.edu (courses.engr.illinois.edu)|130.126.151.9|:80... connected.  
HTTP request sent, awaiting response... 301 Moved Permanently  
Location: https://courses.engr.illinois.edu/cs498aml/sp2019/homeworks/test_file.zip [following]  
--2019-05-02 04:53:56-- https://courses.engr.illinois.edu/cs498aml/sp2019/homeworks/test_file.zip  
Connecting to courses.engr.illinois.edu (courses.engr.illinois.edu)|130.126.151.9|:443... connected.  
HTTP request sent, awaiting response... 200 OK  
Length: 3841776 (3.7M) [application/x-zip-compressed]  
Saving to: 'test_file.zip'  
  
test_file.zip      100%[=====] 3.66M  5.06MB/s    in 0.7s  
  
2019-05-02 04:53:56 (5.06 MB/s) - 'test_file.zip' saved [3841776/3841776]  
  
Archive: test_file.zip  
inflating: test_file.pt
```

Then use your model to predict the label of the test images. Fill the remaining code below, where x has two dimensions (batch_size x one image size). Remember to reshape x accordingly before feeding it into your model. The submission.txt should contain one predicted label (0~9) each line. Submit your submission.txt to the competition in gradscoope.

```
1 import torch.utils.data as Data  
2
```

```
3 test_file = 'test_file.pt'
4 pred_file = 'submission.txt'
5
6 f_pred = open(pred_file, 'w')
7 tensor = torch.load(test_file)
8 torch_dataset = Data.TensorDataset(tensor)
9 test_loader = torch.utils.data.DataLoader(torch_dataset, batch_size, shuffle=False, num_workers=2)
10
11 model.eval()
12 with torch.no_grad():
13     for ele in test_loader:
14         x = ele[0]
15         image = x.reshape(-1,3,32,32)
16         image = image.to(device)
17         preds = model(image)
18
19         pred_labels = torch.argmax(preds, dim=1)
20         for pred_label in pred_labels:
21             f_pred.write(str(pred_label.item()))
22             f_pred.write('\n')
23
24 f_pred.close()
```

▼ Report

Part 0: Imports and Basic Setup (5 Points)

Nothing to report for this part. You will be just scored for finishing the setup.

Part 1: Fully connected neural networks (25 Points)

Test (on validation set) accuracy (5 Points): 0.39

Test loss (5 Points): 1.70

Training time (5 Points): 85.6s

Plots:

- Plot a graph of accuracy on validation set vs training steps (5 Points)
- Plot a graph of loss on validation set vs training steps (5 Points)

Part 2: Convolution Network (Basic) (35 Points)

Tensor dimensions: A good way to debug your network for size mismatches is to print the dimension of output after every layers:

(10 Points)

Output dimension after 1st conv layer: (8,16, 32, 32)

Output dimension after 1st max pooling: (8, 16, 16, 16)

Output dimension after 2nd conv layer: (8,16, 16, 16)

Output dimension after flatten layer: (8, 4096)

Output dimension after 1st fully connected layer: (8,64)

Output dimension after 2nd fully connected layer: (8,10)

Test (on validation set) Accuracy (5 Points): 0.55

Test loss (5 Points): 1.25

Training time (5 Points): 128.4s

Plots:

- Plot a graph of accuracy on validation set vs training steps (5 Points)
- Plot a graph of loss on validation set vs training steps (5 Points)

Part 3: Convolution Network (Add one or more suggested changes) (35 Points)

Describe the additional changes implemented, your intuition for as to why it works, you may also describe other approaches you experimented with (10 Points):

We experimented with various parameters and concluded a final model:

TEST MODEL1: Added 2 additional convolutional layers(total 4) with 16 output channels each because it will extract more subtle features from the image. Added maxpool(2) after each convolutional layer, dropout with probability of 0.2 after the last convolutional layer and batch normalization. Increased epochs to 10. Changed Optimizer from RMSprop to SGD. This gave the test accuracy ~ 64%

TEST MODEL 2: Modified TEST MODEL 1 by changing output channels to 32 in all convolutional layers. This gave the test accuracy ~ 72%

TEST MODEL 3: Modified TEST MODEL 2 dropout rate from 0.2 to 0.4. This gave the test accuracy ~ 71%

FINAL MODEL: Modified TEST MODEL 2 by changing output channels to 128. This gave the test accuracy ~ 79%

Test (on validation set) Accuracy (5 Points): 0.79

Test loss (5 Points): 0.62

Training time (5 Points): 324.5s

Plots:

- Plot a graph of accuracy on validation set vs training steps (5 Points)
- Plot a graph of loss on validation set vs training steps (5 Points)

10 bonus points will be awarded to top 3 scorers on leaderboard (in case of tie for 3rd position everyone tied for 3rd position will get the bonus)

1